

Problem statement

- Implement SGD Classifier with Logloss and L2 regularization Using SGD without using sklearn's SGD function and at the end compare our implementation with the sklearn's implementation...

Importing packages

In [373]:

```
import numpy as np
import pandas as pd
from sklearn.datasets import make_classification
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn import linear_model
import math
```

Creating custom dataset

In [374]:

```
X, y = make_classification(n_samples=50000, n_features=15, n_informative=10, n_redundant=5,
                          n_classes=2, weights=[0.7], class_sep=0.7, random_state=15)
# make_classification is used to create custom dataset
# Please check this Link (https://scikit-learn.org/stable/modules/generated/sklearn.datasets.make\_classification.html) for more d
```

In [375]:

```
X.shape, y.shape
```

Out[375]:

```
((50000, 15), (50000,))
```

Splitting data into train and test

In [376]:

```
# you need not standardize the data as it is already standardized
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=15)
```

In [377]:

```
X_train.shape, y_train.shape, X_test.shape, y_test.shape
```

Out[377]:

```
((37500, 15), (37500,), (12500, 15), (12500,))
```

SGD classifier

In [378]:

```
# alpha : float
# Constant that multiplies the regularization term.

# eta0 : double
# The initial Learning rate for the 'constant', 'invscaling' or 'adaptive' schedules.

clf = linear_model.SGDClassifier(eta0=0.0001, alpha=0.0001, loss='log', random_state=15, penalty='l2', tol=1e-3, verbose=2, learn
clf
# Please check this documentation (https://scikit-learn.org/stable/modules/generated/sklearn.linear\_model.SGDClassifier.html)
```

Out[378]:

```
SGDClassifier(eta0=0.0001, learning_rate='constant', loss='log',
              random_state=15, verbose=2)
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

In [379]:

```
clf.fit(X=X_train, y=y_train) # fitting our model
```

```
-- Epoch 1
Norm: 0.77, NNZs: 15, Bias: -0.316653, T: 37500, Avg. loss: 0.455552
Total training time: 0.01 seconds.
-- Epoch 2
Norm: 0.91, NNZs: 15, Bias: -0.472747, T: 75000, Avg. loss: 0.394686
Total training time: 0.01 seconds.
-- Epoch 3
Norm: 0.98, NNZs: 15, Bias: -0.580082, T: 112500, Avg. loss: 0.385711
Total training time: 0.02 seconds.
-- Epoch 4
Norm: 1.02, NNZs: 15, Bias: -0.658292, T: 150000, Avg. loss: 0.382083
Total training time: 0.03 seconds.
-- Epoch 5
Norm: 1.04, NNZs: 15, Bias: -0.719528, T: 187500, Avg. loss: 0.380486
Total training time: 0.03 seconds.
-- Epoch 6
Norm: 1.05, NNZs: 15, Bias: -0.763409, T: 225000, Avg. loss: 0.379578
Total training time: 0.04 seconds.
-- Epoch 7
Norm: 1.06, NNZs: 15, Bias: -0.795106, T: 262500, Avg. loss: 0.379150
Total training time: 0.04 seconds.
-- Epoch 8
Norm: 1.06, NNZs: 15, Bias: -0.819925, T: 300000, Avg. loss: 0.378856
Total training time: 0.05 seconds.
-- Epoch 9
Norm: 1.07, NNZs: 15, Bias: -0.837805, T: 337500, Avg. loss: 0.378585
Total training time: 0.06 seconds.
-- Epoch 10
Norm: 1.08, NNZs: 15, Bias: -0.853138, T: 375000, Avg. loss: 0.378630
Total training time: 0.06 seconds.
Convergence after 10 epochs took 0.06 seconds
```

C:\Users\natar\anaconda3\envs\tf_gpu\lib\site-packages\sklearn\linear_model_stochastic_gradient.py:173: FutureWarning: The loss 'log' was deprecated in v1.1 and will be removed in version 1.3. Use 'loss='log_loss'' which is equivalent.

```
warnings.warn(
```

Out[379]:

```
SGDClassifier(eta0=0.0001, learning_rate='constant', loss='log',
              random_state=15, verbose=2)
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

In [380]:

```
clf.coef_, clf.coef_.shape, clf.intercept_
#clf.coef_ will return the weights
#clf.coef_.shape will return the shape of weights
#clf.intercept_ will return the intercept term
```

Out[380]:

```
(array([[ -0.42336692,  0.18547565, -0.14859036,  0.34144407, -0.2081867 ,
          0.56016579, -0.45242483, -0.09408813,  0.2092732 ,  0.18084126,
          0.19705191,  0.00421916, -0.0796037 ,  0.33852802,  0.02266721]]),
 (1, 15),
 array([-0.8531383]))
```

Custom implementation of SGD classifier

Initialize weights

In [381]:

```
def initialize_weights(row_vector):
    ''' In this function, we will initialize our weights and bias'''
    # zeros_like function to initialize zero: https://docs.scipy.org/doc/numpy/reference/generated/numpy.zeros_like.html
    w=np.zeros_like(row_vector)
    b=0 #initializing bias to zero
    return w,b
```

In [382]:

```
dim=X_train[0]
w,b = initialize_weights(X_train[0])
print('w =', (w))
print('b =', str(b))
```

```
w = [0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
b = 0
```

Compute sigmoid

$$\text{sigmoid}(z) = 1/(1 + \exp(-z))$$

In [384]:

```
def sigmoid(z):
    ''' In this function, we will return sigmoid of z'''
    if z >= 0:
        z = np.exp(-z)
        return 1 / (1 + z)
    else:
        z = np.exp(z)
        return z / (1 + z)
```

Compute loss

$$\text{logloss} = -1 * \frac{1}{n} \sum_{\text{foreach } Y_t, Y_{\text{pred}}} (Y_t \log_{10}(Y_{\text{pred}}) + (1 - Y_t) \log_{10}(1 - Y_{\text{pred}}))$$

In [386]:

```
def logloss(y_true,y_pred):
    #while dealing with numpy arrays you can use vectorized operations for quicker calculations as compared to using loops
    #https://www.pythonlikeyoumeanit.com/Module3_IntroducingNumpy/VectorizedOperations.html
    #https://www.geeksforgeeks.org/vectorized-operations-in-numpy/

    loss = -np.mean(y_true*(np.log10(y_pred+1e-9)) + (1-y_true)*np.log10(1-y_pred+1e-9)) #to avoid division by zero error add 1e
    return loss
```

Compute gradient w.r.to 'w'

$$dw^{(t)} = x_n(y_n - \sigma((w^{(t)})^T x_n + b^t)) - \frac{\lambda}{N} w^{(t)}$$

In [388]:

```
def gradient_dw(x,y,w,b,alpha,N):
    '''In this function, we will compute the gradient w.r.to w'''
    dw = x*(y-sigmoid(np.dot(w.T,x)+b))*((alpha/N)*w)
    return dw
```

Compute gradient w.r.to 'b'

$$db^{(t)} = y_n - \sigma((w^{(t)})^T x_n + b^t)$$

In [390]:

```
def gradient_db(x,y,w,b):
    '''In this function, we will compute gradient w.r.to b'''
    db = y-sigmoid(np.dot(w,x)+b)
    return db
```

In [392]:

```
# prediction function used to compute predicted_y given the dataset X
def pred(w,b, X):
    N = len(X)
    predict = []
    for i in range(N):
        z=np.dot(w,X[i])+b
        predict.append(sigmoid(z))
    return np.array(predict)
```

Implementing logistic regression

In [393]:

```
def train(X_train,y_train,X_test,y_test,epochs,alpha,eta0):
    ''' In this function, we will implement logistic regression'''

    train_loss = []
    test_loss = []
    w,b = initialize_weights(X_train[0]) # Initialize the weights
    #code to perform SGD
    for i in range(epochs):
        # for every epoch
        for idx,x in enumerate(X_train) :
            # for every data point(X_train,y_train)
            dw=gradient_dw(x,y_train[idx],w,b,alpha,len(X_train)) #computing gradient w.r.to w
            db=gradient_db(x,y_train[idx],w,b) #computing gradient w.r.to b
            w+=eta0*dw #update w #Here eta0 is learning rate
            b+=eta0*db #update b

        y_pred_tr=pred(w,b,X_train) #preding y for the given x_train using logistic function
        tr_loss=logloss(y_train,y_pred_tr) #calculating log loss for train datapoints
        train_loss.append(tr_loss)

        y_pred_te=pred(w,b,X_test) #preding y for the given x_test using logistic function
        te_loss=logloss(y_test,y_pred_te) #calculating log loss for test datapoints
        test_loss.append(te_loss)
        if i>2:
            if train_loss[-1]-train_loss[-2] <= 1e-3: #if the first 3 decimal places of train loss does not change for the co
                break

    return w,b,train_loss,test_loss
```

In [394]:

```
alpha=0.001
eta0=0.001
N=len(X_train)
epochs=20
w,b,train_loss,test_loss=train(X_train,y_train,X_test,y_test,epochs,alpha,eta0)
```

In [395]:

```
#print the value of weights w and bias b
print(w)
print(b)
```

```
[-0.41395135  0.19245248 -0.1500517  0.32635332 -0.22516469  0.58646629
 -0.42720474 -0.10028062  0.2148384  0.15555088  0.17880941 -0.01318672
 -0.0649683  0.36313892 -0.00985027]
-0.9016634051738529
```

In [396]:

```
# these are the results we got after we implemented sgd and found the optimal weights and intercept
w=clf.coef_, b=clf.intercept_
```

Out[396]:

```
(array([[ 0.00941556,  0.00697683, -0.00146134, -0.01509075, -0.01697799,
          0.02630051,  0.02522009, -0.00619249,  0.0055652 , -0.02529038,
          -0.0182425 , -0.01740588,  0.0146354 ,  0.0246109 , -0.03251748]]),
 array([-0.04852511]))
```

Goal of this project

Compare your implementation and SGDClassifier's the weights and intercept, make sure they are as close as possible i.e difference should be in order of 10^{-2}

Grader function - 1

In [397]:

```
#this grader function should return True
#the difference between custom weights and clf.coef_ should be less than or equal to 0.05
def difference_check_grader(w,b,coef,intercept):
    val_array=np.abs(np.array(w-coef))
    assert(np.all(val_array<=0.05))
    print('The custom weights are correct')
    return True
difference_check_grader(w,b,clf.coef_,clf.intercept_)
```

The custom weights are correct

Out[397]:


True

Plot train and test loss vs epochs

plot epoch number on X-axis and loss on Y-axis and make sure that the curve is converging

In [398]:

```
#reference taken from :https://www.oreilly.com/Library/view/machine-Learning-with/9781787121515/cf4693e9-207a-4f74-826c-8190c3c8d
from matplotlib import pyplot as plt
plt.figure(figsize=[8,6])
plt.plot(test_loss,'r',linewidth=1.0)
plt.plot(train_loss,'b',linewidth=1.0)
plt.legend(['Training loss', 'test Loss'],fontsize=9)
plt.xlabel('Epochs',fontsize=16)
plt.ylabel('Loss',fontsize=16)
plt.title('Loss Curves',fontsize=1)
plt.show()
```

the data is well fitted with the model in few epochs itself.we get the **train log loss** of **0.1650592828382934** and **test log loss** of **0.16614668451454206**