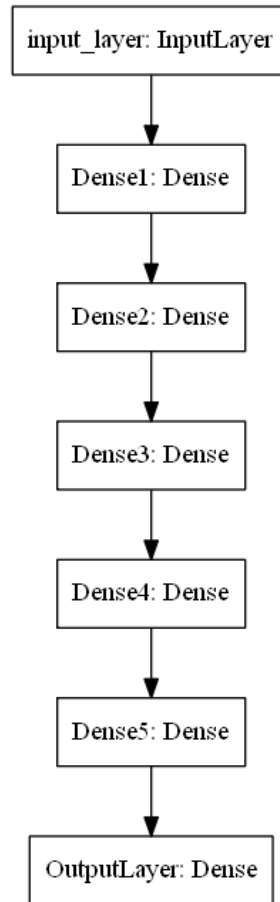


## Problem statement

- Creating custom Call Backs in Tensorflow with the given model...

1. Download the data from [here \(https://drive.google.com/file/d/15dCNcmKskcFVjs7R0EIQkR61Ex53uJpM/view?usp=sharing\)](https://drive.google.com/file/d/15dCNcmKskcFVjs7R0EIQkR61Ex53uJpM/view?usp=sharing). we have to use data.csv file for this project

2. Code the model to classify data like below image. we can use any number of units in your Dense layers.



In [ ]:

```

import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
import tensorflow as tf
from tensorflow.keras.layers import Dense, Input, Activation
from tensorflow.keras.models import Model
import random as rn
from tensorflow import keras
import datetime, os

from keras.callbacks import Callback
from sklearn.metrics import roc_auc_score, f1_score

```

In [ ]:

```

from google.colab import files

uploaded = files.upload()

```

Choose Files No file chosen

Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.

Saving data.csv to data.csv

In [ ]:

```
import io

data = pd.read_csv(io.BytesIO(uploaded['data.csv']))
print(data)
```

```
      f1      f2  label
0  0.450564  1.074305   0.0
1  0.085632  0.967682   0.0
2  0.117326  0.971521   1.0
3  0.982179 -0.380408   0.0
4 -0.720352  0.955850   0.0
...     ...     ...     ...
19995 -0.491252 -0.561558   0.0
19996 -0.813124  0.049423   1.0
19997 -0.010594  0.138790   1.0
19998  0.671827  0.804306   0.0
19999 -0.854865 -0.588826   0.0
```

[20000 rows x 3 columns]

In [ ]:

```
y=data['label'].values
x=data[['f1','f2']].values
```

In [ ]:

```
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.25, stratify=y)
```

### 3. Writing Callbacks

#### we are implementing the following callbacks

- Write our own callback function, that has to print the micro F1 score and AUC score after each epoch. we can not use `tf.keras.metrics` for calculating AUC and F1 score.
- Save your model at every epoch if your validation accuracy is improved from previous epoch.
- we have to decay learning based on below conditions
  - Cond1. If your validation accuracy at that epoch is less than previous epoch accuracy, we have to decrease the learning rate by 10%.
  - Cond2. For every 3rd epoch, decay our learning rate by 5%.
- If we are getting any NaN values (either weights or loss) while training, we have to terminate your training.
- we have to stop the training if our validation accuracy is not increased in last 2 epochs.
- Use tensorboard for every model and analyse our scalar plots and histograms.

In [ ]:

```
from tensorflow.keras.callbacks import ModelCheckpoint
from tensorflow.keras.callbacks import EarlyStopping
from tensorflow.keras.callbacks import ReduceLROnPlateau
from tensorflow.keras.callbacks import LearningRateScheduler
```

In [ ]:

```
# Load the TensorBoard notebook extension
%load_ext tensorboard
```

In [ ]:

```

class custom_metrics(tf.keras.callbacks.Callback):
    def __init__(self, validation_data, patience=2):
        self.X_val, self.y_val = validation_data
        self.patience = patience
        self.best_val_acc = 0.0
        self.best_model_weights = None

    def on_train_begin(self, logs={}):
        self.learning_rate = self.model.optimizer.learning_rate

    def on_epoch_end(self, epoch, logs={}):
        # Calculate the micro F1 score and AUC score
        y_pred = (np.asarray(self.model.predict(self.X_val))).round()
        #y_pred = np.argmax(self.model.predict(self.X_val), axis=-1)
        f1 = f1_score(self.y_val, y_pred, average='micro')
        auc_score = roc_auc_score(self.y_val, y_pred)
        print(f'Micro F1 Score: {f1:.4f}, AUC Score: {auc_score:.4f}')

        #If you are getting any NaN values(either weights or loss) while training, you have to terminate your training.
        loss = logs.get('loss')
        model_weights = self.model.get_weights()
        if loss is not None and model_weights is not None:
            if np.isnan(loss) or np.isinf(loss) or np.any([np.any(np.isnan(x)) for x in model_weights]): #updated line to check
                print("Invalid loss and terminated at epoch {}".format(epoch))
                self.model.stop_training = True

        # Decay the Learning rate based on the validation accuracy
        val_acc = logs.get('val_accuracy')
        if val_acc < self.best_val_acc:
            self.learning_rate.assign(self.learning_rate * 0.9)
            self.model.optimizer.learning_rate = self.learning_rate
        if epoch % 3 == 0:
            self.learning_rate.assign(self.learning_rate * 0.95)
            self.model.optimizer.learning_rate = self.learning_rate

```

In [ ]:

```

#custom metric calling
cus_metrics=custom_metrics(validation_data=(x_test,y_test))

```

In [ ]:

```

#save best model
filepath="model_save/weights-{epoch:02d}-{val_accuracy:.4f}.hdf5"
checkpoint = ModelCheckpoint(filepath=filepath, monitor='val_accuracy', verbose=1, save_best_only=True, mode='auto')

```

In [ ]:

```

#stop the training if your validation accuracy is not increased in last 2 epochs.
earlystop = EarlyStopping(monitor='val_accuracy', min_delta=0.00001, patience=2, verbose=1)

```

In [ ]:

```

#tensorboard
logdir = os.path.join("logs", datetime.datetime.now().strftime("%Y%m%d-%H%M%S"))
tensorboard_callback = tf.keras.callbacks.TensorBoard(logdir, histogram_freq=1)

```

**Model-1**

1. Use tanh as an activation for every layer except output layer.
2. use SGD with momentum as optimizer.
3. use RandomUniform(0,1) as initializer.
3. Analyze your output and training process.

In [ ]:

```
def create_model_1():
    return tf.keras.models.Sequential([
        tf.keras.layers.Dense(2, activation="tanh", input_shape=(2, ), kernel_initializer=keras.initializers.RandomUniform(minval=-0, maxval=1)),
        tf.keras.layers.Dense(10, activation="tanh", kernel_initializer=keras.initializers.RandomUniform(minval=-0, maxval=1)),
        tf.keras.layers.Dense(10, activation="tanh", kernel_initializer=keras.initializers.RandomUniform(minval=-0, maxval=1)),
        tf.keras.layers.Dense(10, activation="tanh", kernel_initializer=keras.initializers.RandomUniform(minval=-0, maxval=1)),
        tf.keras.layers.Dense(10, activation="tanh", kernel_initializer=keras.initializers.RandomUniform(minval=-0, maxval=1)),
        tf.keras.layers.Dense(1, activation='softmax', kernel_initializer=keras.initializers.RandomUniform(minval=-0, maxval=1))
    ])

```

In [ ]:

```
model_1=create_model_1()
optimizer=tf.keras.optimizers.SGD(learning_rate=0.01, momentum=0.9, nesterov=False, name='SGD') #SGD
model_1.compile(optimizer,
                loss='BinaryCrossentropy',
                metrics=['accuracy'])
model_1.fit(x=x_train,
            y=y_train,
            epochs=15,
            validation_data=(x_test, y_test), callbacks=[cus_metrics, checkpoint, tensorboard_callback, earlystop ])

```

Epoch 1/15

1/469 [.....] - ETA: 23:34 - loss: 1.3704 - accuracy: 0.4375

WARNING:tensorflow:Callback method `on\_train\_batch\_end` is slow compared to the batch time (batch time: 0.0025s vs `on\_train\_batch\_end` time: 0.0031s). Check your callbacks.

157/157 [=====] - 0s 1ms/step

Micro F1 Score: 0.5000, AUC Score: 0.5000

Epoch 1: val\_accuracy improved from -inf to 0.50000, saving model to model\_save/weights-01-0.5000.hdf5

469/469 [=====] - 5s 5ms/step - loss: 0.7246 - accuracy: 0.5000 - val\_loss: 0.6934 - val\_accuracy: 0.5000

Epoch 2/15

157/157 [=====] - 0s 1ms/step

Micro F1 Score: 0.5000, AUC Score: 0.5000

Epoch 2: val\_accuracy did not improve from 0.50000

469/469 [=====] - 2s 5ms/step - loss: 0.6953 - accuracy: 0.5000 - val\_loss: 0.6962 - val\_accuracy: 0.5000

Epoch 3/15

157/157 [=====] - 0s 1ms/step

Micro F1 Score: 0.5000, AUC Score: 0.5000

Epoch 3: val\_accuracy did not improve from 0.50000

469/469 [=====] - 2s 4ms/step - loss: 0.6954 - accuracy: 0.5000 - val\_loss: 0.6933 - val\_accuracy: 0.5000

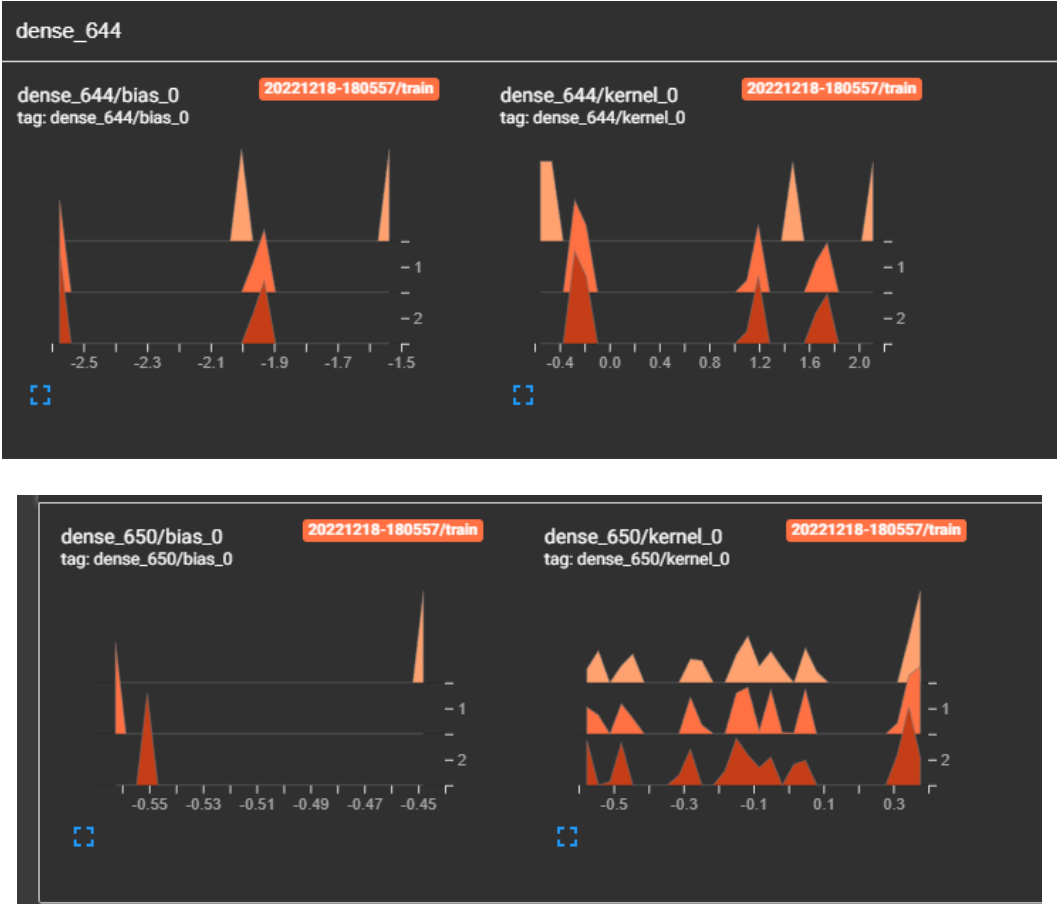
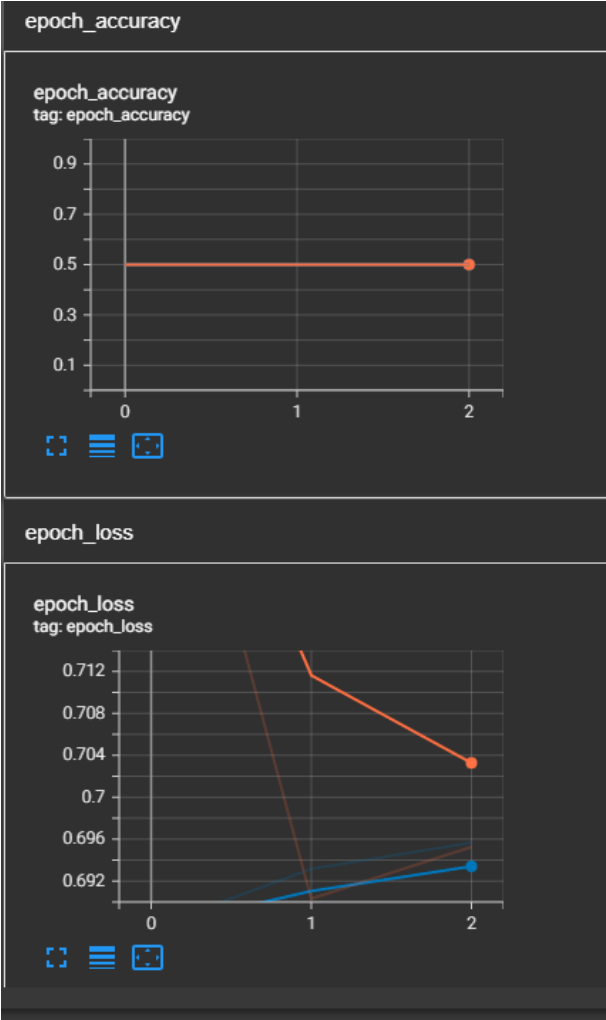
Epoch 3: early stopping

Out[14]:

&lt;keras.callbacks.History at 0x7f0ee669e790&gt;

In [ ]:

!rm -rf ./logs/



Observations

- from the above plot 1 we can see that the accuracy did not changed much through all the epoch
- from the plot 2(dense\_644) we can see that wights in the first dense layer are distributed in the range between -0.4 and 2 and it does not change much during the all epoch..since derivative of tanh function's max value is 1, it somewhat avoid the vanishing gradient problem but still weights are not distributed well...
- from the plot 3(dense\_650) we can see that initial wights in the final dense layer are distributed in the range between -0.5 and 0.3 and it does not change much during the all epoch and since this is last layer the vanishing gradient does not affect this layers weight much so that the weights are reasonably distributed well...

**Model-2**

1. Use relu as an activation for every layer except output layer.
2. use SGD with momentum as optimizer.
3. use RandomUniform(0,1) as initilizer.
3. Analyze your output and training process.

In [ ]:

```
def create_model_2():
    return tf.keras.models.Sequential([
        tf.keras.layers.Dense(2,activation="relu",input_shape=(2,)),kernel_initializer=keras.initializers.RandomUniform(minval=-0, maxval=1)),
        tf.keras.layers.Dense(10, activation="relu",kernel_initializer=keras.initializers.RandomUniform(minval=-0, maxval=1)),
        tf.keras.layers.Dense(10, activation="relu",kernel_initializer=keras.initializers.RandomUniform(minval=-0, maxval=1)),
        tf.keras.layers.Dense(10, activation="relu",kernel_initializer=keras.initializers.RandomUniform(minval=-0, maxval=1)),
        tf.keras.layers.Dense(10, activation="relu",kernel_initializer=keras.initializers.RandomUniform(minval=-0, maxval=1)),
        tf.keras.layers.Dense(1, activation='softmax',kernel_initializer=keras.initializers.RandomUniform(minval=-0, maxval=1))
    ])

```

In [ ]:

```
model_2=create_model_2()
optimizer=tf.keras.optimizers.SGD(learning_rate=0.01, momentum=0.9, nesterov=False, name='SGD')
model_2.compile(optimizer,
                loss='BinaryCrossentropy',
                metrics=['accuracy'])
model_2.fit(x=x_train,
            y=y_train,
            epochs=15,
            validation_data=(x_test, y_test),callbacks=[cus_metrics,checkpoint,tensorboard_callback,earlystop ])

```

Epoch 1/15

1/469 [.....] - ETA: 3:51 - loss: 195.2360 - accuracy: 0.6562

WARNING:tensorflow:Callback method `on\_train\_batch\_end` is slow compared to the batch time (batch time: 0.0023s vs `on\_train\_batch\_end` time: 0.0024s). Check your callbacks.

157/157 [=====] - 0s 1ms/step

Micro F1 Score: 0.5000, AUC Score: 0.5000

Epoch 1: val\_accuracy did not improve from 0.50000

469/469 [=====] - 3s 5ms/step - loss: 1.1083 - accuracy: 0.5000 - val\_loss: 0.6935 - val\_accuracy: 0.5000

Epoch 2/15

157/157 [=====] - 0s 1ms/step

Micro F1 Score: 0.5000, AUC Score: 0.5000

Epoch 2: val\_accuracy did not improve from 0.50000

469/469 [=====] - 2s 5ms/step - loss: 0.6934 - accuracy: 0.5000 - val\_loss: 0.6932 - val\_accuracy: 0.5000

Epoch 3/15

157/157 [=====] - 0s 1ms/step

Micro F1 Score: 0.5000, AUC Score: 0.5000

Epoch 3: val\_accuracy did not improve from 0.50000

469/469 [=====] - 2s 5ms/step - loss: 0.6934 - accuracy: 0.5000 - val\_loss: 0.6932 - val\_accuracy: 0.5000

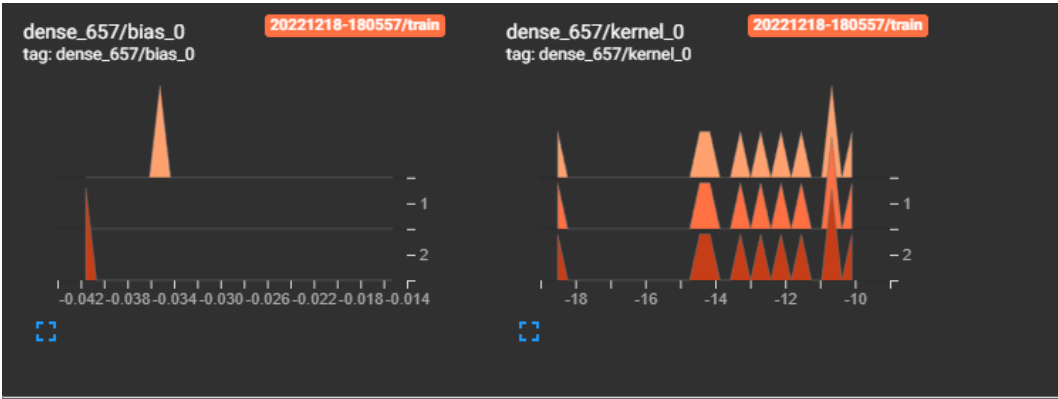
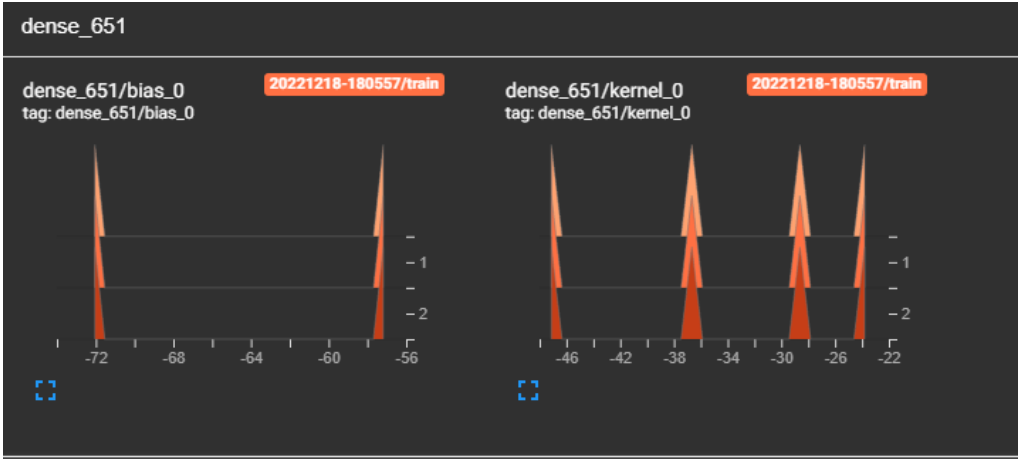
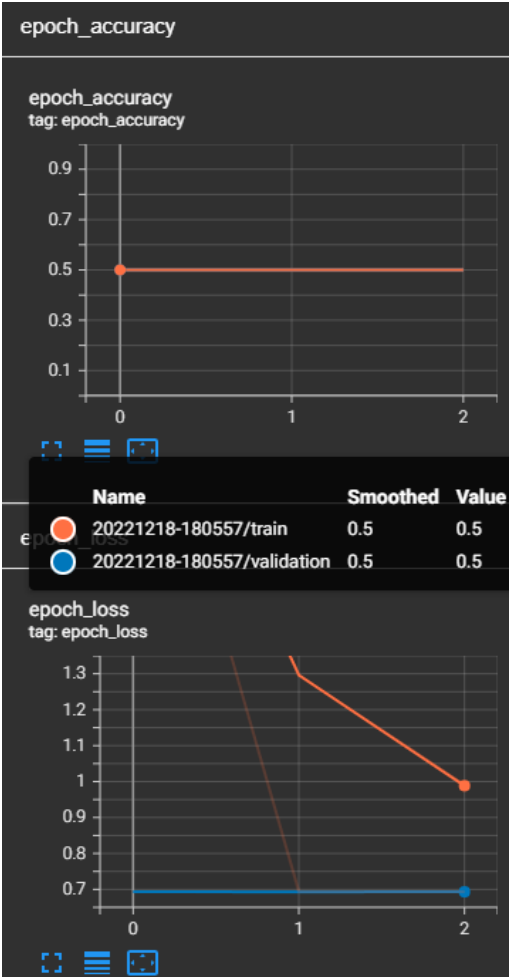
Epoch 3: early stopping

Out[20]:

&lt;keras.callbacks.History at 0x7f0fa5e0e8b0&gt;

In [ ]:

```
!rm -rf ./logs/
```



## Observations

- from the above plot 1 we can see that the accuracy did not changed much through all the epoch and train loss reduces from 2.3 to 1.
- from the plot 2(dense\_651) we can see that wights in the first dense layer are distributed in the range between -46 and -22 and it does not change much during the all epoch..since derivative of relu function 1 for the positive value, it avoid the vanishing gradient problem but there is a high chance for exploding gradient problem...and due to that weights in the first layers are not well distributed
- from the plot 3(dense\_657) we can see that wights in the final dense layer are distributed in the range between -18 and -0 and it does not change much during the all epoch and since this is last layer there is no vanishing gradient probelm as well as not much of emploting gradient probelm also and the weights are reasonably distributed well...

### Model-3

1. Use relu as an activation for every layer except output layer.
2. use SGD with momentum as optimizer.
3. use he\_uniform() as initilizer.
3. Analyze your output and training process.

In [ ]:

```
def create_model_3():
    return tf.keras.models.Sequential([
        tf.keras.layers.Dense(2,activation="relu",input_shape=(2,),kernel_initializer=keras.initializers.he_uniform()), #relu
        tf.keras.layers.Dense(10, activation="relu",kernel_initializer=keras.initializers.he_uniform()), #he_uniform()
        tf.keras.layers.Dense(10, activation="relu",kernel_initializer=keras.initializers.he_uniform()),
        tf.keras.layers.Dense(10, activation="relu",kernel_initializer=keras.initializers.he_uniform()),
        tf.keras.layers.Dense(10, activation="relu",kernel_initializer=keras.initializers.he_uniform()),
        tf.keras.layers.Dense(1, activation='sigmoid',kernel_initializer=keras.initializers.he_uniform())
    ])
```



In [ ]:

```

model_3=create_model_3()
optimizer=tf.keras.optimizers.SGD(learning_rate=0.01, momentum=0.9, nesterov=False, name='SGD')
model_3.compile(optimizer,
                 loss='BinaryCrossentropy',
                 metrics=['accuracy'])
model_3.fit(x=x_train,
            y=y_train,
            epochs=10,
            validation_data=(x_test, y_test),callbacks=[cus_metrics,checkpoint,tensorboard_callback,earlystop ])

```

Epoch 1/10

157/157 [=====] - 0s 1ms/step

Micro F1 Score: 0.5718, AUC Score: 0.5718

Epoch 1: val\_accuracy improved from 0.50000 to 0.57180, saving model to model\_save/weights-01-0.5718.hdf5

469/469 [=====] - 3s 5ms/step - loss: 0.6804 - accuracy: 0.5577 - val\_loss: 0.6756 - val\_accuracy: 0.5718

Epoch 2/10

157/157 [=====] - 0s 1ms/step

Micro F1 Score: 0.5842, AUC Score: 0.5842

Epoch 2: val\_accuracy improved from 0.57180 to 0.58420, saving model to model\_save/weights-02-0.5842.hdf5

469/469 [=====] - 2s 5ms/step - loss: 0.6692 - accuracy: 0.5848 - val\_loss: 0.6688 - val\_accuracy: 0.5842

Epoch 3/10

157/157 [=====] - 0s 1ms/step

Micro F1 Score: 0.5912, AUC Score: 0.5912

Epoch 3: val\_accuracy improved from 0.58420 to 0.59120, saving model to model\_save/weights-03-0.5912.hdf5

469/469 [=====] - 2s 4ms/step - loss: 0.6676 - accuracy: 0.5893 - val\_loss: 0.6657 - val\_accuracy: 0.5912

Epoch 4/10

157/157 [=====] - 0s 1ms/step

Micro F1 Score: 0.5982, AUC Score: 0.5982

Epoch 4: val\_accuracy improved from 0.59120 to 0.59820, saving model to model\_save/weights-04-0.5982.hdf5

469/469 [=====] - 2s 5ms/step - loss: 0.6647 - accuracy: 0.5927 - val\_loss: 0.6661 - val\_accuracy: 0.5982

Epoch 5/10

157/157 [=====] - 0s 1ms/step

Micro F1 Score: 0.5864, AUC Score: 0.5864

Epoch 5: val\_accuracy did not improve from 0.59820

469/469 [=====] - 2s 5ms/step - loss: 0.6628 - accuracy: 0.5917 - val\_loss: 0.6695 - val\_accuracy: 0.5864

Epoch 6/10

157/157 [=====] - 0s 2ms/step

Micro F1 Score: 0.5982, AUC Score: 0.5982

Epoch 6: val\_accuracy did not improve from 0.59820

469/469 [=====] - 2s 5ms/step - loss: 0.6619 - accuracy: 0.5985 - val\_loss: 0.6663 - val\_accuracy: 0.5982

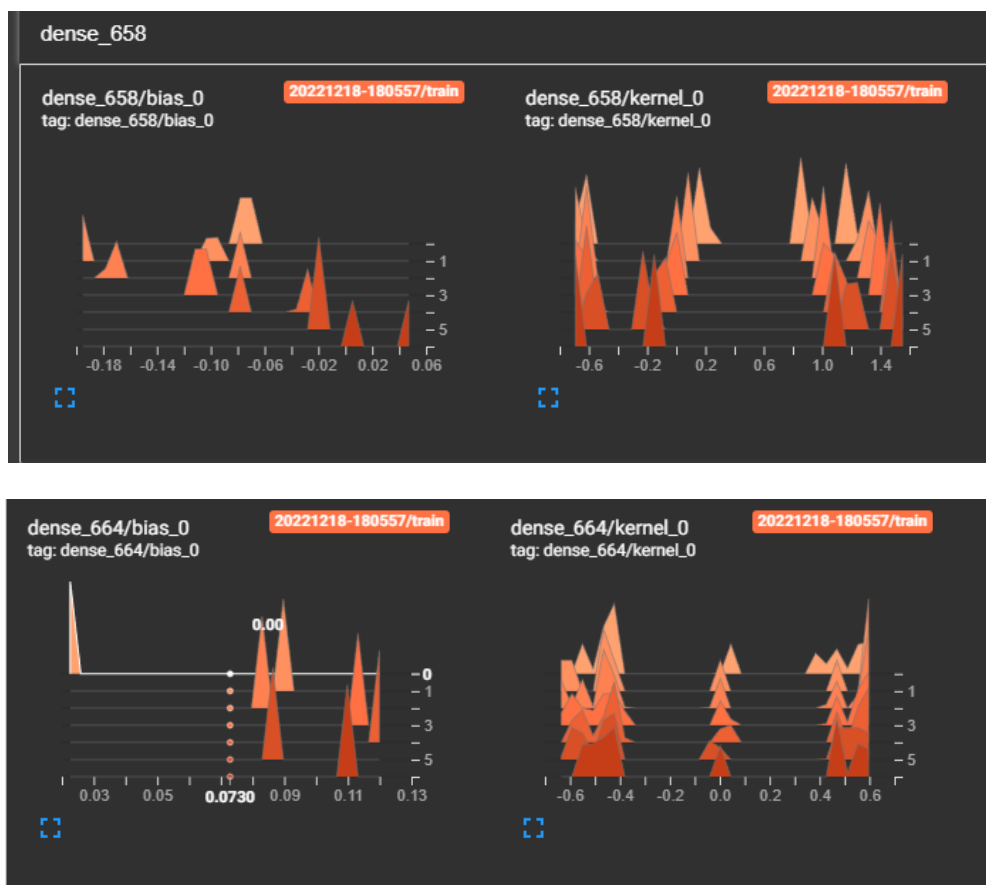
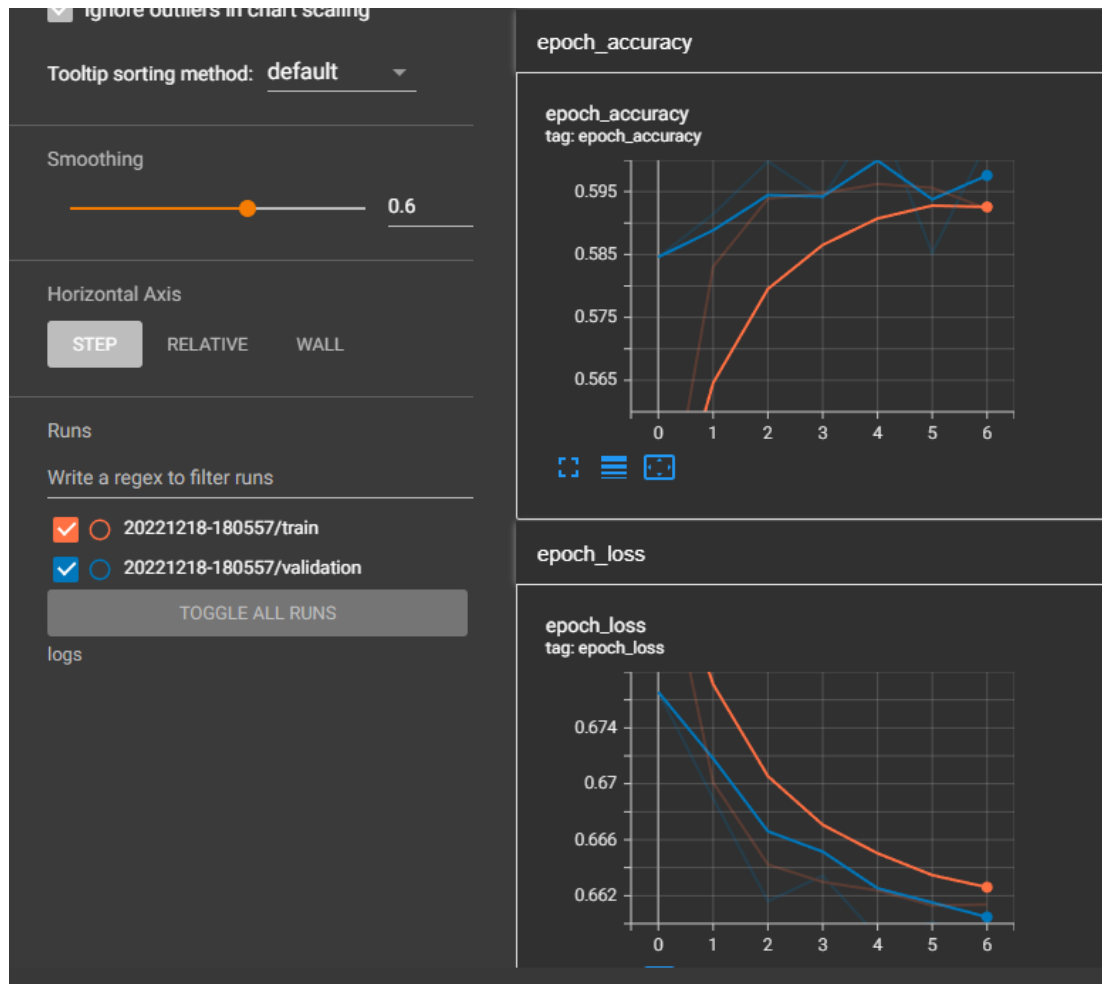
Epoch 6: early stopping

Out[24]:

&lt;keras.callbacks.History at 0x7f0fa5791d60&gt;

In [ ]:

!rm -rf ./logs/



## Observations

- from the above plot 1 we can see that the accuracy did not changed much through all the epoch and train loss and test loss reduces from 0.68 to 0.66.

- from the plot 2(dense\_658) we can see that he\_uniform() initialisation helps wights in the first dense layer are distributed in the range between -0.6 and +1.4..so there are distributed around the mean and it does change much during the all epoch..since derivative of relu function 1 for the positive value, it avoid the vanishing gradient problem but there is a high chance for exploding gradient problem..but this method of weight intialisation avoid those exploding gradient problem also....but weights in the first layers are not well distributed
- from the plot 3(dense\_664) we can see that wights in the final dense layer are distributed in the range between -0.6 and 0.6 and it does not change much during the all epoch and since this is last layer there is no vanishing gradient probelm as well as not much of emploting gradient probelm also but the weights are not reasonably distributed well...

#### Model-4

1. Try with any values to get better accuracy/f1 score.

In [ ]:

*#model 4 build with relu as a activation function and glorot\_normal() as the weight intialisation method and adam as an optimizer*

```
def create_model_4():
    return tf.keras.models.Sequential([
        tf.keras.layers.Dense(10,activation="relu",input_shape=(2,),kernel_initializer=keras.initializers.glorot_normal()),
        tf.keras.layers.Dropout(0.8),
        tf.keras.layers.Dense(10, activation="relu",kernel_initializer=keras.initializers.glorot_normal()),
        tf.keras.layers.Dropout(0.8),
        tf.keras.layers.Dense(10, activation="relu",kernel_initializer=keras.initializers.glorot_normal()),
        tf.keras.layers.Dropout(0.8),
        tf.keras.layers.Dense(10, activation="relu",kernel_initializer=keras.initializers.glorot_normal()),
        tf.keras.layers.Dropout(0.8),
        tf.keras.layers.Dense(10, activation="relu",kernel_initializer=keras.initializers.glorot_normal()),
        tf.keras.layers.Dropout(0.8),
        tf.keras.layers.Dense(1, activation='sigmoid',kernel_initializer=keras.initializers.glorot_normal())
    ])
```

In [ ]:

```
model_4=create_model_4()
optimizer=tf.keras.optimizers.SGD(learning_rate=0.01, momentum=0.9, nesterov=False, name='adam')
#optimizer = tf.keras.optimizers.Adam(0.01)
model_4.compile(optimizer,
                loss='BinaryCrossentropy',
                metrics=['accuracy'])
model_4.fit(x=x_train,
            y=y_train,
            epochs=5,
            validation_data=(x_test, y_test),callbacks=[cus_metrics,checkpoint,tensorboard_callback,earlystop ]
            )
```

Epoch 1/5

1/469 [.....] - ETA: 4:42 - loss: 0.6784 - accuracy: 0.4688

WARNING:tensorflow:Callback method `on\_train\_batch\_end` is slow compared to the batch time (batch time: 0.0025s vs `on\_train\_batch\_end` time: 0.0027s). Check your callbacks.

157/157 [=====] - 0s 1ms/step

Micro F1 Score: 0.5000, AUC Score: 0.5000

Epoch 1: val\_accuracy did not improve from 0.59820

469/469 [=====] - 3s 5ms/step - loss: 0.7130 - accuracy: 0.4981 - val\_loss: 0.6932 - val\_accuracy: 0.5000

Epoch 2/5

157/157 [=====] - 0s 1ms/step

Micro F1 Score: 0.5000, AUC Score: 0.5000

Epoch 2: val\_accuracy did not improve from 0.59820

469/469 [=====] - 2s 5ms/step - loss: 0.6933 - accuracy: 0.5021 - val\_loss: 0.6934 - val\_accuracy: 0.5000

Epoch 3/5

157/157 [=====] - 0s 1ms/step

Micro F1 Score: 0.5000, AUC Score: 0.5000

Epoch 3: val\_accuracy did not improve from 0.59820

469/469 [=====] - 2s 5ms/step - loss: 0.6934 - accuracy: 0.4948 - val\_loss: 0.6933 - val\_accuracy: 0.5000

Epoch 3: early stopping

Out[28]:

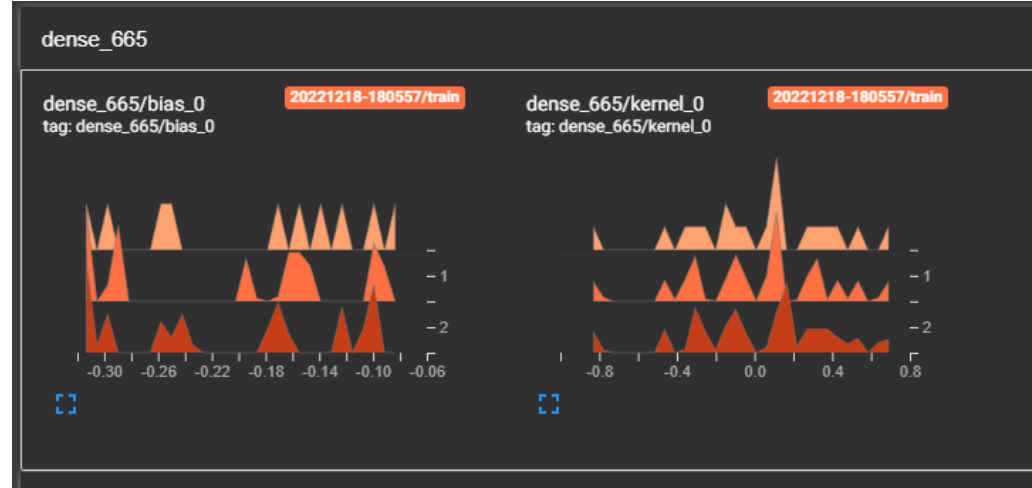
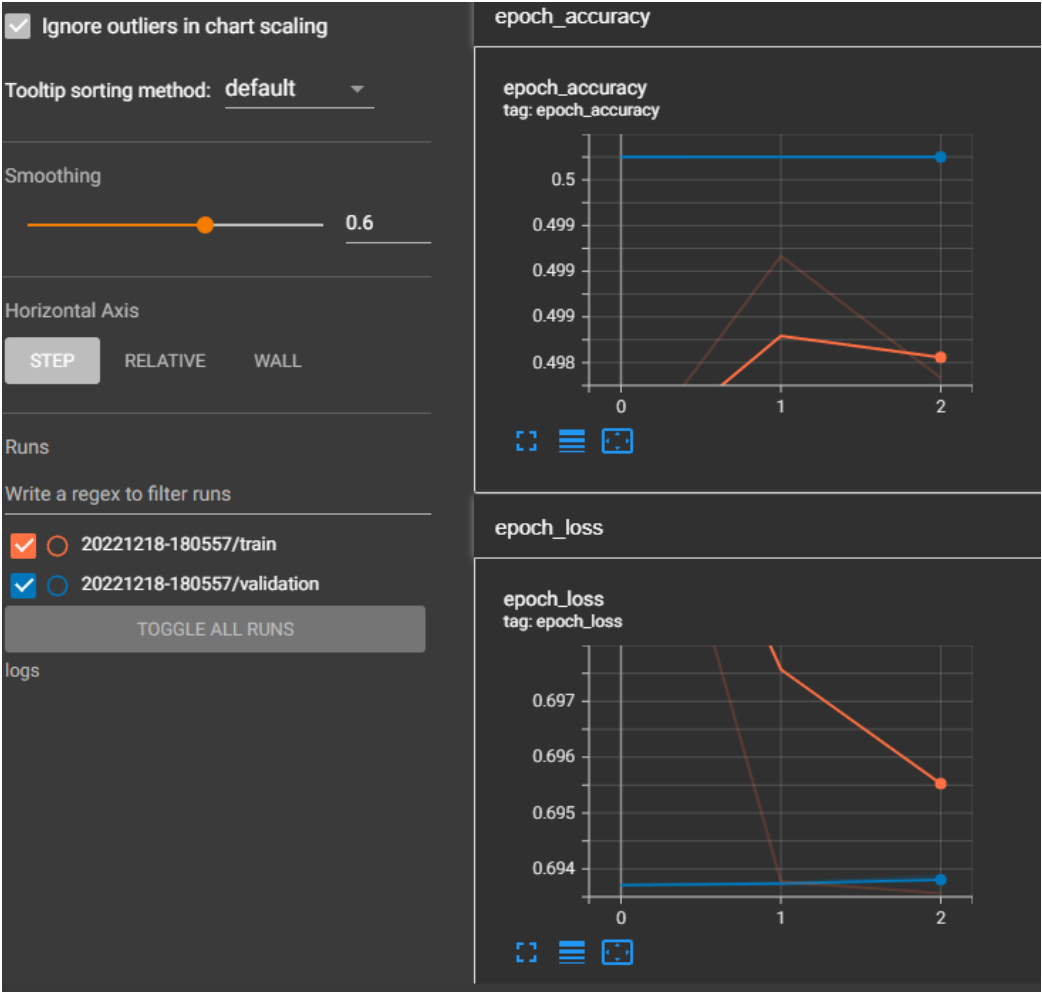
<keras.callbacks.History at 0x7f0eccf17c70>


In [ ]:

```
!kill 2158
```

In [ ]:

```
!rm -rf ./logs/
```





## Observations

- from the above plot 1 we can see that the accuracy did not changed much through all the epoch and train loss and test loss is not change significantly
- from the plot 2(dense\_665) we can see that glorot\_normal() initialisation helps wights in the first dense layer are distributed in the range between -0.8 and 0.8..and the weights are well distributed...so there are distributed around the mean and it does not change much during the all epoch..since derivative of relu function 1 for the positive value, it avoid the vanishing gradient problem but there is a high chance for exploding gradient problem..but this method of weight intialisation avoid those exploding gradient problem also...
- from the plot 3(dense\_671) we can see that wights in the final dense layer are distributed in the range between -0.5 and 0.7 and it does not change much during the all epoch and since this is last layer there is no vanishing gradient probelm as well as not much of emploting gradient probelm also and the weights are reasonably distributed well...

In [ ]: